

线性回归分析 boston 数据集

```
In [1]: import numpy as np
```

```
In [2]: import pandas as pd
```

```
In [3]: from sklearn.datasets import load_boston
```

```
In [4]: from sklearn.cross_validation import train_test_split
```

```
In [5]: from sklearn.linear_model import LinearRegression
```

```
In [6]: from sklearn.metrics import mean_squared_error
```

```
In [7]: import matplotlib.pyplot as plt
```

```
In [8]: from sklearn.preprocessing import PolynomialFeatures
```

将数据集划分成训练集,开发集,测试集

```
In [9]: data = load_boston()
```

```
In [10]: X = data['data']
```

```
In [11]: Y = data['target']
```

```
In [13]: X_train,X_test_all,y_train,y_test_all =  
         train_test_split(X,Y,test_size=0.3,random_state=9)
```

```
In [14]: X_dev,X_test,y_dev,y_test =  
         train_test_split(X_test_all,y_test_all,  
                           test_size=0.3,random_state=9)
```

建模

```
In [15]: model = LinearRegression(normalize=True,fit_intercept=True)
```

```
In [16]: model.fit(X_train,y_train)
```

```
In [17]: y_predicted = model.predict(X_train)
```

绘制残差

```
In [18]: %matplotlib
```

```
In [19]: plt.cla()
```

```
In [20]: plt.xlabel("Predicted y")
```


```
In [21]: plt.ylabel("Residual")

In [22]: plt.title("Residual Plot")

In [23]: plt.figure(1)

In [24]: diff = y_train - y_predicted

In [25]: plt.plot(y_predicted,diff,'go')
```



查看模型系数

```
In [26]: for i,coef in enumerate(model.coef_):
...:     print "Coefficient %d %0.3f" % (i+1,coef)
...:
Coefficient 1 -0.109
Coefficient 2 0.043
Coefficient 3 0.053
Coefficient 4 2.237
```



```
In [38]: model_poly.fit(X_train_poly,y_train)

In [39]: y_predicted = model_poly.predict(X_train_poly)

In [40]: print "MSE = %0.2f"
          % (mean_squared_error(y_train,y_predicted))
MSE = 5.46    # 很有效

# 将模型作用到开发集上
In [41]: y_predicted = model_poly.predict(X_dev_poly)

In [42]: print "MSE = %0.2f"
          % (mean_squared_error(y_dev,y_predicted))
MSE = 13.23

# 将模型应用到测试集
In [43]: X_test_poly = poly_features.transform(X_test)

In [44]: y_predicted = model_poly.predict(X_test_poly)

In [45]: print "MSE = %0.2f"
          % (mean_squared_error(y_test,y_predicted))
MSE = 15.13

In [46]: y_predicted = model.predict(X_test)

In [47]: print "MSE = %0.2f"
          % (mean_squared_error(y_test,y_predicted))
MSE = 21.66
```